Politecnico di Milano

A.A. 2016–2017

Software Engineering 2: "PowerEnJoy"

Code Inspection Document

Pietro Ferretti, Nicole Gervasoni, Danilo Labanca

February 5, 2017

Contents

1 Assigned Class

```
apache-ofbiz-16.11.01/framework/service/
      src/main/java/org/apache/ofbiz/service/job/JobManager.java
```

2 Functional Role

2.1 Service Engine

Introduction

Services are independent pieces of logic which when placed together process many different types of business requirements. Services can be of many different types: Workflow, Rules, Java, SOAP, BeanShell, etc. A service with the type Java is much like an event where it is a static method, however with the Services Framework we do not limit to web based applications. Services require input parameters to be in a Map and the results are returned in a Map as well. This is nice since a Map can be serialized and stored or passed via HTTP (SOAP).

Services are defined through the Service Definition and are assigned to a specific Service Engine. Each Service Engine is responsible for invoking the defined service in an appropriate way. Since services are not tied to web based applications this allows services to run when there is no response object available. This allows services to be scheduled to run at specific times to run in the background via the Job Scheduler.

Services have the ability to call other services. So, chaining small services together to accomplish a larger task makes reusing existing services much easier.

Services which are declared in a component (service-resource in ofbiz-component.xml file) are reachable from anywhere in OFBiz, and even outside using the export feature. Though this possibility is not commonly used (no examples OOTB), it's also possible to create services which are specific to an application: restricted to be available only in that application. For that, you put the service definition and implementation files under the WEB-INF directory. You can also override a service by using the same name down in the deployment context (which is first framework, then themes, then applications, then specialpurpose, then hot-deploy). This is handy, but beware if not wanted...

When used in a web application services are available to web events, which allow events to stay small and reuse existing logic in the Services Framework. Also, services can be defined as 'exportable' which means they are allowed to be accessed by outside parties. Currently there is a SOAP EventHandler which allows services to be made available via SOAP. Other forms of remote invocation may be added to the framework in the future.

2.2 Package: Job etc.

Job Scheduler

The overhauled job scheduler is now integrated with the services framework. This is the most appropriate place for a scheduler. Since it cannot be guarantee

that an HttpServletRequest and HttpServletResponse object will be available when a job is ready to run, it does not make sense to integrate with the web controller. Plus, this feature is most useful when not limited to web environments.

The scheduler is a multi-threaded component with a single thread used for job managing/scheduling and separate threads used for invocation of each service. When a job is scheduled to run, the scheduler will call the service dispatcher associated with the job to invoke the service in its own thread. This will prevent long or time consuming jobs from slowing down other jobs in the queue.

The scheduler now supports the iCalendar rule structure for recurrence. The jobs are no longer stored in an XML file and each is part of a ServiceDispatcher. There is one Job Scheduler for each ServiceDispatcher (which means there is only one per GenericDelegator as well).

How it works:

The best usage example of the scheduler is an asynchronous service call. When an asynchronous service is invoked, it is passed to the Job Scheduler to be queued to run. A recurrence entry is created (RecurrenceInfo and RecurrenceRule entities are created), the job is stored, (JobSandbox entity is created) and the context (Map) is serialized and stored (RuntimeData entity is created). The scheduler then adds the job to the top of the list of scheduled jobs (asynchronous services do not have any delay time) and invoked.

Jobs are no longer defined in an XML file. This has been moved to the JobSandbox entity. There is a web based client in planning for adding predefined jobs to the queue, but currently the entities will have to be created by hand.

2.3   Class: Job Manager

```
/**
 * Job manager. The job manager queues and manages jobs. Client
     code can queue a job to be run immediately
 * by calling the runJob({@link #runJob(Job)}) method, or schedule
     a job to be run later by calling the
 * {@link #schedule(String, String, String, Map, long, int, int,
     int, long, int)} method.
 * Scheduled jobs are persisted in the JobSandbox entity.
 * <p>A scheduled job's start time is an approximation - the
     actual start time will depend
 * on the job manager/job poller configuration (poll interval) and
     the load on the server.
 * Scheduled jobs might be rescheduled if the server is busy.
     Therefore, applications
 * requiring a precise job start time should use a different
     mechanism to schedule the job.</p>
 */
```

The job manager queues and manages jobs.

Client code can queue a job to be run immediately by calling the runJob method, or schedule a job to be run later by calling the schedule method.

Scheduled jobs are persisted in the JobSandbox entity. A scheduled job's start time is an approximation - the actual start time will depend on the job manager/job poller configuration (poll interval) and the load on the server. Scheduled jobs might be rescheduled if the server is busy. Therefore, applications requiring a precise job start time should use a different mechanism to schedule the job.

3 List of Issues

3.1 Naming Conventions

1. Meaningful Names sono tutti buoni

2. One-character variables Ok. There are no one-character variables.

3. Class names Ok. Every class name is in mixed case and properly capitalized.

4. Interface names Ok. No interfaces are declared. (se ce ne sono) every interface used by the code is in mixed case and properly capitalized.

5. Method names Ok. Every method name is a verb. Every method name is camelCase and properly capitalized.

6. Class variables Ok. Every class variable is in mixed case and properly capitalized.

7. Constants NO. module and istanceId are immutable, so they can be considered constant. They should be capitalized. registeredManagers is fine because it's mutable

3.2 Indentation

8. Number of spaces Ok. The code is consistently indented with 4 spaces.

9. No tabs for indentation Ok. No tabs are used to indent the code.

3.3 Braces

10. Consistent bracing style Ok. The code is consistently braced following the *"Kernighan and Ritchie"* style.

11. One-line statements bracing NO "if" riga 326, 351, 354

3.4 File Organization

12. Blank lines as separation Ok. Blank lines are present between each method, around imports and variable declarations. Most of the methods also begin with a Javadoc.

13. Where practical, line length under 80 characters NOPE righe 73, 74, 89, 126, 147, 150, 154, 156, 161, 182, 186-190, 195, 198, 201, A great number of lines exceed 80 characters

14. Line length always under 120 characters NEPPURE righe 74, 186, 198, 217, 221, 222, 261-264, 273, 311, 315, 317, 387, 409, 429, 453, 498, 543, 560, 561 le dichiarazioni dei metodi sono lunghissime e wrappate poco

3.5 Wrapping Lines

15. Line breaks after commas and operators NO riga 152, la virgola dovrebbe stare sopra

16. Higher-level breaks are used Ok. Non ci sono line break con operatori

17. Statements are aligned to previous ones Ok. Sì, per tutti

3.6 Comments

3.7 Java Source Files

20. Single public class or interface Ok. Job manager is the only public class declared in the file. There are no other classes.

21. The public class is the first class in the file Ok. Job manager is the only public class declared in the file. There are no other classes.

22. External program interfaces are consistent with the Javadoc Ok abbiamo vari metodi pubblici: getter: - getDelegator - getDispatcher - getInstance - getPoolState poi altre robe - isAvailable - reloadCrashedJobs - runJob - schedule di tutti i tipi
    la Javadoc parla di runJob e schedule

23. The Javadoc is complete NO.

- No javadoc for 'module'! line 71

- No javadoc for 'instanceId'! line 71

- No javadoc for reloadCrashedJob!! line 304

- Missing @return tag on getInstance, line 88

- Missing @return tag on getDelegator, line 119

- Missing @return tag on getDispatcher, line 124

- Missing @param tag for 'limit' on poll, line 174

- Missing @return tag on poll, line 174

- Missing @param tag for 'job' on runJob, line 363

- Missing @throws tag for 'JobManagerException' on runJob, line 363

- Missing @throws tag for 'JobManagerException' on schedule, line 386, 408, 428, 453, 469, 498, 543

assertIsRunning, getRunPools sono private quindi non hanno necessaria-mente bisogno di javadoc

3.8   Package and Import Statements

24. Package statements are first, import statements second Ok. One package statements. All import statements immediately follow.

3.9   Class and Interface Declarations

25. The class declarations should follow a specific order   - javadoc ok - class declaration ok - altri commenti / - static variables ok - public ok - private ok - normal variables - constructors - methods

no, abbiamo variabili statiche, poi un po' di metodi statici, poi variabili normali, poi costruttori (getInstance è un costruttore), setter e getter poi un metodo statico (ma private!!)

26. Methods are grouped by functionality   Ok
assertIsRunning getInstance shutDown
getDelegator getDispatcher getPoolState
isAvailable getRunPools pool reloadCrashedJobs runJob schedule

27. The code is free of duplicates, long methods, big classes, breaking encapsulation, and coupling and cohesion are adequate small class duplicates? no short methods no breaking encapsulation

low/loose coupling -> ci sono un sacco di delegator e dispatcher high cohesion -> tutti i metodi servono a runnare/queuare jobs

3.10 Initialization and Declarations

3.11 Method Calls

3.12 Arrays

No off-by-one errors in array indexing Ok. The only indexing is with foreach, no off-by-one errors.

No out-of-bounds indexes Ok. No number indexing.

Constructors are called when a new array item is desired Ok. quali nuovi array? non ce ne sono

3.13 Object Comparison

Objects are compared with equals Ok. There are no object comparisons.

3.14 Output Format

Displayed output is free of spelling and grammatical errors riga 156: Debug.logWarning(e, "Exception thrown while check lock on JobManager : " + instanceId, module); dovrebbe essere "while checking"

riga 182: Debug.logWarning("Unable to locate DispatchContext object; not running job!", module); dovrebbe essere "job:", come negli altri log di debug

Error messages are comprehensive and useful Sì

Output is formatted correctly in terms of line breaks and spacing No line breaks in outputs Some debug outputs don't have a trailing space

3.15 Computation, Comparisons and Assignments

3.16 Exceptions

3.17 Flow of Control

All switch cases are addressed with a break Ok, no switch statements.

All switch statements have a default branch Ok, no switch statements.

All loops are correctly formed, with appropriate initialization, increments and termination expressions Ok. All for loops are foreach, everything is fine. The while loop at line 219: GenericValue jobValue = jobsIterator.next(); while (jobValue != null) { jobValue = jobsIterator.next(); tutto ok, l'iteratore va avanti finché non finiscono i valori, poi esce dal while

while a riga 275 uguale

a posto

4 Other Problems

alla riga 576: viene chiamato new Long(0), che è meno efficiente di valueOf(0) Using new Integer(int) is guaranteed to always result in a new object whereas Integer.valueOf(int) allows caching of values to be done by the compiler, class library, or JVM. Using of cached values avoids object allocation and the code will be faster. Values between -128 and 127 are guaranteed to have corresponding cached instances and using valueOf is approximately 3.5 times faster than using constructor. For values outside the constant range the performance of both styles is the same

se ci sono tante richieste sostituirlo migliorerebbe le prestazioni ma comunque problema minore, basso impatto

5 Effort Spent

- Pietro Ferretti: hours of work

- Nicole Gervasoni: hours of work

- Danilo Labanca: hours of work

6 Revisions

6.1 Changelog

- CID v1.0, published on February 5, 2017